

HW2_starter_code

February 10, 2025

1 Starter Code for ECE 580 HW2

Demo code for solving LASSO:

$$\arg \min_w \|Y - Xw\|_2^2 + \alpha \|w\|_1$$

```
[36]: import torch
      # lasso code
      from sklearn.linear_model import Lasso

      torch.manual_seed(0)
      X = torch.randn(4,6)    # 4 samples, 6 features
      Y = torch.randn(4)      # Target variable

      # Set the regularization strength alpha = 10
      myalpha = 0.01

      # Fit the Lasso model
      lasso = Lasso(alpha=myalpha, fit_intercept=False, max_iter=10000)
      lasso.fit(X.numpy(), Y.numpy())

      # Display the coefficients
      print("Target:", Y)
      print("X:", X)
      print("w:", lasso.coef_)
      print("\033[91mNote: if alpha is unspecified, sklearn.linear_model.Lasso_
      ↪defaults to alpha=1.0 \033[0m")
```

```
Target: tensor([0.3704, 1.4565, 0.9398, 0.7748])
X: tensor([[ -1.1258, -1.1524, -0.2506, -0.4339,  0.8487,  0.6920],
          [-0.3160, -2.1152,  0.4681, -0.1577,  1.4437,  0.2660],
          [ 0.1665,  0.8744, -0.1435, -0.1116,  0.9318,  1.2590],
          [ 2.0050,  0.0537,  0.6181, -0.4128, -0.8411, -2.3160]])
w: [ 0.6800222 -0.12440585  0.          -0.3782028  0.9276224 -0.          ]
Note: if alpha is unspecified, sklearn.linear_model.Lasso defaults to
alpha=1.0
```

Possibly handy function for rasterizing images: torch.reshape()

```
[37]: torch.manual_seed(0)
v = torch.randn(4,8)
print('initial shape of v: ', v.shape)
v_vec = v.reshape(32)
print('after reshaping v: ', v_vec.shape)
print('verify that v==v.reshape(4,8).reshape(32): ', (v-v_vec.reshape(4,8)).
      ↪norm().item())
```

```
initial shape of v: torch.Size([4, 8])
after reshaping v: torch.Size([32])
verify that v==v.reshape(4,8).reshape(32): 0.0
```

Code for loading image

```
[38]: import torchvision.transforms as transforms
from PIL import Image
import matplotlib.pyplot as plt
# Load the BMP image using PIL
image_path = "fishing_boat.bmp" # Replace with your BMP file path
original_image = Image.open(image_path)

# Convert to grayscale using torchvision transforms
transform = transforms.Compose([
    transforms.Grayscale(num_output_channels=1), # Convert to 1-channel
    ↪grayscale
    transforms.ToTensor() # Convert to a PyTorch tensor
])
# Apply the transform
grayscale_image = transform(original_image)[0,:,:]

# Check the shape of the grayscale image tensor
print("Grayscale Image Shape:", grayscale_image.shape)

plt.imshow(grayscale_image, cmap="gray") # Use cmap="gray" for grayscale
    ↪display
plt.title("Grayscale fishing_boat")
plt.axis("off") # Remove axis for better visualization
plt.show()

# Save the grayscale image as a BMP for verification (optional)
save_path = "grayscale_fishing_boat.bmp"
transforms.ToPILImage()(grayscale_image).save(save_path)
print(f"Grayscale image saved to {save_path}")
```

```
Grayscale Image Shape: torch.Size([200, 192])
```

Grayscale fishing_boat



Grayscale image saved to grayscale_fishing_boat.bmp

2 Question 1

2.1 a).

```
[39]: # Recall from HW1, our proposed model included adding a constant term to the
      ↪ input

      # load dataste
      X_train, y_train, _, _ = torch.load('../HW1/cleaned_automobile_train_dataset.
      ↪ pt').values()
      X_test, y_test, _, _ = torch.load('../HW1/cleaned_automobile_test_dataset.pt').
      ↪ values()

      # Add a constant term to the input
      X_train = torch.cat((X_train, torch.ones(X_train.shape[0], 1)), dim=1).cpu().
      ↪ numpy()
      X_test = torch.cat((X_test, torch.ones(X_test.shape[0], 1)), dim=1).cpu().
      ↪ numpy()
```

```

from sklearn.model_selection import cross_validate
from sklearn.linear_model import LinearRegression

model = LinearRegression()

# Perform 6-fold cross-validation
results = cross_validate(model, X_train, y_train, cv=6,
    ↳scoring='neg_mean_squared_error')

mses = results['test_score'] * -1
average_mses = mses.mean()

print(f'{mses=}')
print(f'{average_mses=}')

```

```

mses=array([ 9677476. , 16784012. , 7996756.5, 7787531.5, 10094987. ,
            15623169. ])
average_mses=np.float64(11327322.0)

```

Reported MSES: 9677476, 16894012, 7996756.5, 7787531.5, 10094987, 15623169 Reported Average MSE: 11327322

We see that there are quite some discrepancies depending the choice of validation sets, especially in the scenario of limited data.

2.2 b).

```

[ ]: # first compute regular linear regression
linear = LinearRegression()
linear.fit(X_train, y_train)
beta = linear.coef_

# find the l2norm of beta
l2norm = torch.tensor(beta).norm()
norm_low = 0.1 * l2norm
norm_high = 0.5 * l2norm

from sklearn.linear_model import Ridge
import numpy as np
# select lambda in a logarithmic scale
norms = []
for i in 50 / 2**np.linspace(0, 9, 10):
    ridge = Ridge()
    ridge.fit(X_train, y_train)
    beta = ridge.coef_
    l2norm = torch.tensor(beta).norm()
    if norm_low <= l2norm <= norm_high:
        print(f'{i=}, {l2norm=}, {norm_low=}, {norm_high=}')


```

```
norms.append(( , l2norm))
```

```
[42]: from matplotlib import pyplot as plt
```

```
lams, l2norms = zip(*norms)
plt.plot(lams, l2norms)
```

```
[42]: [<matplotlib.lines.Line2D at 0x7f3e8993b9d0>]
```




2.3 c).

```
[ ]: res_train = []
res_test = []
for lam in lams:
    # Perform 6-fold cross-validation
    ridge = Ridge(lam)
    results = cross_validate(ridge, X_train, y_train, cv=6,
    ↪scoring='neg_mean_squared_error', return_train_score=True)
    train_mses = results['train_score'] * -1
```

```
test_mses = results['test_score'] * -1
res_train.append(train_mses.mean())
res_test.append(test_mses.mean())
```

```
[44]: plt.plot(lams, np.log10(res_train), label='train')
plt.plot(lams, np.log10(res_test), label='validation')
plt.legend()
```

```
[44]: <matplotlib.legend.Legend at 0x7f3e89970e10>
```



We see that as lambda increases, although the training mse increases, the validation mse actually decreases, closing the gap between training and validation sets.

2.4 d).

```
[45]: lam_star = 50

# this choice made based on validation error is different from the one made
↳ based on the train error.
ridge = Ridge(lam_star)
```

```

ridge.fit(X_train, y_train)
y_hat = ridge.predict(X_test)
if y_test.device != 'cpu':
    y_test = y_test.cpu().numpy()
mse = ((y_hat - y_test)**2).mean()

print(f'test_{mse=}')

# 6-fold cross-validation
from sklearn.model_selection import cross_val_score
cross_score = cross_val_score(ridge, X_train, y_train, cv=6,
    ↳scoring='neg_mean_squared_error')
cross_mse = -cross_score
print(f'{cross_mse=}')

```

```

test_mse=np.float32(194545970.0)
cross_mse=array([ 9994327.,  9460999.,  7936231.,  8400108.,  9928843.,
16531553.])

```

2.5 e).

```

[46]: # pick largest lambda
lams = 3000000 / 2**np.linspace(0, 9, 10)

if y_train.device != 'cpu':
    y_train = y_train.cpu().numpy()

for lam in lams:
    lasso = Lasso(lam)
    lasso.fit(X_train, y_train)
    w = np.linalg.norm(lasso.coef_)
    print(w)

```

```

0.0
5.7737513
8.676704
10.128181
11.009962
11.5499115
11.829032
46.624863
68.739136
80.1849

```

```

[47]: # s = 3000000 / 2**np.linspace(0, 9, 10)
from sklearn.linear_model import Lasso

```


```

ws = []
zeros = []
for lam in lams:
    lasso = Lasso(lam)
    # cross validation
    cross_result = cross_validate(lasso, X_train, y_train, cv=6,
    ↪scoring='neg_mean_squared_error', return_estimator=True)
    # cross_mses = -cross_result['test_score']
    # print(f'{cross_mses.mean()}')
    # mses.append(cross_mses.mean())
    # estimators = cross_result['estimator']
    estimators = cross_result['estimator']
    w = [np.linalg.norm(est.coef_, 1) for est in estimators]
    zero = [np.sum(est.coef_ == 0) for est in estimators]
    ws.append(np.mean(w))
    zeros.append(np.mean(zero))

plt.plot(np.log(lams), ws, label='L1Norm')
plt.plot(np.log(lams), zeros, label='0 Entries')
plt.xlabel('log()')
plt.legend()

```

[47]: <matplotlib.legend.Legend at 0x7f3e898d5810>

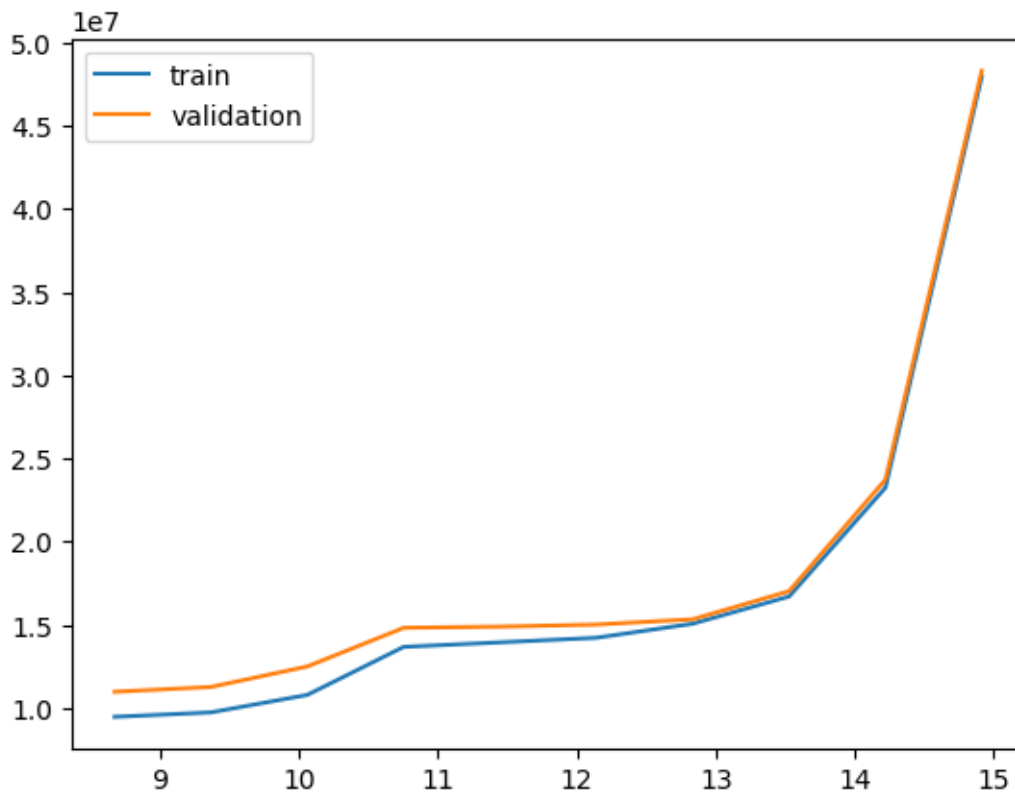


2.6 f).

```
[48]: train_mses = []
      test_mses = []
      for lam in lams:
          lasso = Lasso(lam)
          results = cross_validate(lasso, X_train, y_train, cv=6,
          ↪scoring='neg_mean_squared_error', return_estimator=True,
          ↪return_train_score=True)
          train_mse = -results['train_score']
          test_mse = -results['test_score']
          train_mses.append(train_mse.mean())
          test_mses.append(test_mse.mean())

      plt.plot(np.log(lams), train_mses, label='train')
      plt.plot(np.log(lams), test_mses, label='validation')
      plt.legend()
```

[48]: <matplotlib.legend.Legend at 0x7f3e8970d1d0>



2.7 g).

the 4th lambda (187500) was chosen because while it has some descent train/validation accuracy, it had very small gap between train and validation dataset

```
[49]: # best lambda
lam_star = lams[4]
print(f'{lam_star=}')

lasso = Lasso(lam_star)
lasso.fit(X_train, y_train)
y_hat = lasso.predict(X_test)
mse = ((y_hat - y_test)**2).mean()
print(f'{mse=}')

# cross validation mse
cross_score = cross_val_score(lasso, X_train, y_train, cv=6,
                               scoring='neg_mean_squared_error')
cross_mse = -cross_score
```

```
print(f'{cross_mse=}')
```

```
lam_star=np.float64(187500.0)
mse=np.float32(162735650.0)
cross_mse=array([10961815., 12712907., 7754417., 17780134., 9201014.,
31744704.])
```

3 Problem 2

3.1 1 a).

```
[50]: import torchvision.transforms as transforms
      from PIL import Image
```

```
[51]: signal = np.array([1, 2, 3, 4])
      N = 4
      coefficients = np.zeros(N)
      for mu in range(N):
          alpha = np.sqrt(1/N) if mu == 0 else np.sqrt(2/N)
          for x in range(N):
              coefficients[mu] += signal[x] * np.cos(np.pi * (2*x + 1) * mu / (2*N))
          coefficients[mu] *= alpha

      print(coefficients)
```

```
[ 5.00000000e+00 -2.23044250e+00 -6.28036983e-16 -1.58512668e-01]
```

3.2 1 b).

```
[52]: # def dft_basis(signal, D):
      #     N = signal.shape[0]
      #     coefficients = np.zeros((D, N))
      #     for mu in range(D):
      #         alpha = np.sqrt(1/N) if mu == 0 else np.sqrt(2/N)
      #         for x in range(N):
      #             coefficients[mu, x] = signal[x] * np.cos(np.pi * (2*x + 1) * mu /
      # ↪ (2*N))
      #         coefficients[mu] *= alpha
      #     return coefficients

      signal = np.array([1, 2, 3, 4])
      D = 8
      X = np.linspace(0, 2*np.pi, 100).reshape(-1, 1)
      basis = np.zeros((D, X.shape[0]))
```

```

for mu in range(D):
    alpha = np.sqrt(1/D) if mu == 0 else np.sqrt(2/D)
    for i in range(X.shape[0]):
        x = X[i]
        basis[mu, i] = alpha * np.cos(np.pi * (2*x + 1) * mu / (2*D))
    plt.plot(X, basis[mu], label=f'mu={mu}')

basis = basis.T

# plt.plot(X, basis)
plt.legend()

```


/tmp/ipykernel_2522/3911806360.py:20: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)

```

basis[mu, i] = alpha * np.cos(np.pi * (2*x + 1) * mu / (2*D))

```

[52]: <matplotlib.legend.Legend at 0x7f3e8af82210>



3.3 1 c).

3.4 2 a).

```
[53]: coefficients = np.array([14, -3.5, 3.5, 0])
      g1 = np.zeros(coefficients.shape[0])
      N = coefficients.shape[0]
      for x in range(coefficients.shape[0]):
          for mu in range(N):
              alpha = np.sqrt(1/N) if mu == 0 else np.sqrt(2/N)
              g1[x] += (coefficients[mu] * np.cos(np.pi * (2*x + 1) * mu / (2*N))) * alpha
      g1
```

```
[53]: array([ 6.46351481,  4.30290682,  6.19709318, 11.03648519])
```

3.5 2 b).

```
[54]: coefficients = np.array([14, -3.5, 0, 0])
      g2 = np.zeros(coefficients.shape[0])
      N = coefficients.shape[0]
      for x in range(coefficients.shape[0]):
          for mu in range(N):
              alpha = np.sqrt(1/N) if mu == 0 else np.sqrt(2/N)
              g2[x] += (coefficients[mu] * np.cos(np.pi * (2*x + 1) * mu / (2*N))) * alpha
      g2
```

```
[55]: g2
```

```
[55]: array([4.71351481, 6.05290682, 7.94709318, 9.28648519])
```

3.6 2 c).

```
[56]: g = np.array([4, 6, 8, 10])
      mse1 = ((g - g1)**2).mean()
      mse2 = ((g - g2)**2).mean()
      print(f'{mse1=}')
      print(f'{mse2=}')

```

```
mse1=np.float64(3.318451259143179)
mse2=np.float64(0.2559512591431786)
```

The reconstruction on the Sparse DCT coefficient actually yielded a smaller MSE


3.7 3 a).

```
[57]: # case 1
u, v = 3, 4
P, Q = 8, 8
alpha = np.sqrt(1/P) if u == 1 else np.sqrt(2/P)
beta = np.sqrt(1/Q) if v == 1 else np.sqrt(2/Q)
X = np.linspace(0, 2*np.pi, 100).reshape(-1, 1)
basis1 = alpha * beta * np.cos(np.pi * (2*X - 1) * (u-1) / (2*P)) * np.cos(np.
    ↪pi * (2*X - 1) * (v-1) / (2*Q))
plt.plot(X, basis1, label='basis1')

# case 2
u, v = 5, 2
alpha = np.sqrt(1/P) if u == 1 else np.sqrt(2/P)
beta = np.sqrt(1/Q) if v == 1 else np.sqrt(2/Q)
basis2 = alpha * beta * np.cos(np.pi * (2*X - 1) * (u-1) / (2*P)) * np.cos(np.
    ↪pi * (2*X - 1) * (v-1) / (2*Q))
plt.plot(X, basis2, label='basis2')

plt.legend()
```

[57]: <matplotlib.legend.Legend at 0x7f3e8ae48410>



3.8 3 b).

```
[58]: G = np.array([
        [1, 2],
        [3, 4]
    ])

    # 1. display all basis functions
    T = np.zeros((G.shape[0], G.shape[1]))
    for x in range(1, G.shape[0]+1):
        for y in range(1, G.shape[1]+1):
            for u in range(1, G.shape[0]+1):
                for v in range(1, G.shape[1]+1):
                    alpha = np.sqrt(1/G.shape[0]) if u == 1 else np.sqrt(2/G.
↪shape[0])
                    beta = np.sqrt(1/G.shape[1]) if v == 1 else np.sqrt(2/G.
↪shape[1])
                    # basis = alpha * beta * np.cos(np.pi * (2*x - 1) * (u-1) /
↪(2*G.shape[0])) * np.cos(np.pi * (2*y - 1) * (v-1) / (2*G.shape[1]))
                    basis = alpha * beta * np.cos(np.pi * (2*x - 1) * (u-1) / (2*G.
↪shape[0])) * np.cos(np.pi * (2*y - 1) * (v-1) / (2*G.shape[1]))
                    T[x-1, y-1] += basis

                    # plt.plot(X, basis, label=f'basis_{u}_{v}')

    # plt.legend()
    T
```

```
[58]: array([[2.00000000e+00, 2.22044605e-16],
        [2.22044605e-16, 0.00000000e+00]])
```

```
[59]: # 2. compute the reconstruction of G

    reconstruction = np.zeros(G.shape)
    for x in range(G.shape[0]):
        for y in range(G.shape[1]):
            for u in range(1, G.shape[0]+1):
                for v in range(1, G.shape[1]+1):
                    alpha = np.sqrt(1/G.shape[0]) if u == 1 else np.sqrt(2/G.
↪shape[0])
                    beta = np.sqrt(1/G.shape[1]) if v == 1 else np.sqrt(2/G.
↪shape[1])
                    reconstruction[x,y] += (alpha * beta * np.cos(np.pi * (2*x - 1)
↪* (u-1) / (2*G.shape[0])) * np.cos(np.pi * (2*y - 1) * (v-1) / (2*G.
↪shape[1]))) * G[u-1][v-1]
```

```
reconstruction
```

```
[59]: array([[5., 5.],  
           [5., 5.]])
```

3.9 3 c).


4 Problem 3

4.1 1 a).

```
[60]: from PIL import Image  
  
img = Image.open('fishing_boat.bmp')
```

```
[61]: from matplotlib import pyplot as plt  
from torchvision.transforms import transforms  
# grayscale  
  
transform = transforms.Compose([  
    transforms.Grayscale(num_output_channels=1), # Convert to 1-channel  
    ↪grayscale  
    transforms.ToTensor() # Convert to a PyTorch tensor  
)  
  
grayscale_image = transform(img)[0,:,:]  
  
plt.imshow(grayscale_image, cmap='gray')
```


```
[61]: <matplotlib.image.AxesImage at 0x7f3e8aea3390>
```

4.2 b).

```
[63]: plt.imshow( grayscale_image, cmap='gray')
```


```
[63]: <matplotlib.image.AxesImage at 0x7f3e8adaec10>
```



```
[64]: x = 8*3
      y = 8*6
      block = grayscale_image[x:x+8, y:y+8]
      block.shape

      plt.imshow(block, cmap='gray')
```

```
[64]: <matplotlib.image.AxesImage at 0x7f3e8ac2a850>
```





```
[65]: from numpy import random
      random.seed(10)


      for s in [50, 30, 10]:

          corrupted_block = block.clone().flatten()
          # corrupted_block[torch.randperm(64)[:64-s]] = 0.5
          pixels = torch.randperm(64)[:64-s]
          masks = torch.zeros(64)
          masks[pixels] = 1
          masks = masks.reshape(8, 8)

          plt.imshow(block, cmap='gray')
          for i in range(8):
              for j in range(8):
                  if masks[i, j] == 1:
                      plt.text(j, i, 'X', color='red')
          plt.title(f'S={s}')
          plt.show()
```







4.3 1 d).

4.4 2 a).

```
[73]: # get corrupted image
S = 30
cur = 0

c_block = block.clone()
f_block = c_block.flatten()
pixels = torch.randperm(64)[:64-S]
masks = torch.zeros(64)
masks[pixels] = 1

P, Q = 64, 64
T = np.zeros((P, Q))
# apply dct
for u in range(P):
    for v in range(Q):
        alpha = np.sqrt(1/(P*Q)) if u == 0 else np.sqrt(2/(P*Q))
        T[u, v] = alpha * np.cos(np.pi * (2*v + 1) * u / (2*P*Q))
```

```
# T.shape, T
```

```
[74]: print('Note: the lambda was chosen to be smaller than asked in the assignment,
        ↳because the original lambda yielded all-0 vector')
lam = 0.01
lasso = Lasso(lam)
T1 = T[masks==0]
# T2 = T1[:, masks==0]
corrupted_block = c_block.flatten()[masks == 0]
lasso.fit(T1, corrupted_block.reshape(-1, 1))
print(lasso.coef_)
```

Note: the lambda was chosen to be smaller than asked in the assignment because the original lambda yielded all-0 vector

```
[ 0.  0.  0.  0.  0.  0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0.
 -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0.
 -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0.
 -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0. -0.]
```

```
[75]: # lasso to estimate alpha

from sklearn.linear_model import Lasso
lam = 0.0000003

lasso = Lasso(lam, max_iter=2000000)
lasso.fit(T1, corrupted_block.reshape(-1, 1))
lasso.coef_
```

```
[75]: array([-0.          , -0.          , -0.          , -0.          , -0.          ,
        -0.          , -0.          , -0.          , -0.          , -0.          ,
        -0.          , -0.          , -0.          , -0.          ,  0.          ,
         0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
         0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
         0.          ,  0.          ,  0.          ,  0.          ,  0.24677243,
        0.2906345 , 0.17686363, 0.07700953,  0.          ,  0.          ,
         0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
         0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
         0.          ,  0.          ,  0.          , -0.          , -0.          ,
        -0.          , -0.          , -0.          , -1.11822324])
```

```
[76]: max_mags = np.argsort(-np.abs(lasso.coef_), axis=0, )[:5]
# to u,v index
for pos in max_mags:
    print(f'u,v = {pos/8},{pos%8}, weight={lasso.coef_[pos]}')
```

```

u,v = 7.875,7, weight=-1.1182232371341951
u,v = 5.0,0, weight=0.29063450408137853
u,v = 4.875,7, weight=0.24677242953078446
u,v = 5.125,1, weight=0.17686362879903064
u,v = 5.25,2, weight=0.07700953068137688

```

4.5 2 b).

```

[77]: # reconstruct image

reconstruction = np.zeros(64)


coefs = lasso.coef_

for i in range(64):
    for j in range(64):
        alpha = np.sqrt(1/(64)) if i == 0 else np.sqrt(2/(64))
        reconstruction[i] += alpha * coefs[j] * np.cos(np.pi * (2*j + 1) * i /
↪(2*64))

reconstruction = reconstruction.reshape(8, 8)
plt.imshow(reconstruction, cmap='gray')

```

[77]: <matplotlib.image.AxesImage at 0x7f3e8aad0410>



4.6 2 c).

```
[78]: if block.device != 'cpu':  
        block = block.cpu()  
        reconstruction = torch.tensor(reconstruction).to(block.device)  
mse = ((reconstruction - block)**2).mean()  
print(f'{mse=}')
```

```
mse=tensor(0.5666, dtype=torch.float64)
```

4.7 2 d).

```
[ ]:
```